

01.112 Machine Learning

Yeoh Chuen Wen 1001775 | Christabella Adelina 1001620 | Aditya Manikashetti 1001819
Design Project

Part 2: Simple Sentiment Analysis

I. Approach

The simple sentiment analysis produces the most probable tag sequence for a given token sequence by choosing the highest emission probability of a tag for each given token. Unrecognised words in the input file (a word that appears less than $k=3$ times, the learning threshold) will use the emission probability for the general token #UNK#.

For every input token, our algorithm checks if the token is found in our emission probability table. If it is found in the emission probability table, the tag which produces the highest emission probability for this token will be recorded. If the token is not found in the table, the tag which produces the highest emission probability for the unknown token #UNK# is used instead. The predicted tag for each token is compiled into a list resulting in a predicted tag sequence for a given token sequence.

$$y^* = \arg \max_y e(x|y)$$

The simple method takes the input file tokens and the emission parameters and outputs as a predicted tag sequence (for a given token sequence).

II. Result

	EN	FR	CN	SG
<i>Entity in gold data</i>	226	223	362	1382
<i>Entity in prediction</i>	1201	1149	3318	6542
<i>Correct entity</i>	165	182	183	780
<i>Entity precision</i>	0.1374	0.1584	0.0552	0.1192
<i>Entity recall</i>	0.7301	0.8161	0.5055	0.5644
<i>Entity F</i>	0.2313	0.2653	0.0995	0.1969
<i>Correct sentiment</i>	71	68	57	311
<i>Sentiment precision</i>	0.0591	0.0592	0.0172	0.0475
<i>Sentiment recall</i>	0.3142	0.3049	0.1575	0.2250
<i>Sentiment F</i>	0.0995	0.0991	0.0310	0.0785

Part 3: Viterbi Algorithm

I. Approach

The Viterbi algorithm produces the most probable tag sequence for a given token sequence. This is done by choosing the tag sequence (calculation of the transition parameters) with the highest probability of producing the given token sequence (calculation of the emission parameters). Similarly, unrecognised words in the input file (words that appear less than $k=3$ times, the learning threshold) will use the emission probability for the general token #UNK#.

$$r(y_1, \dots, y_k) = \prod_{i=1}^k a_{y_{i-1}, y_i} \prod_{i=1}^k b_{y_i}(x_i)$$

$$p(x_1, \dots, x_n, y_0, y_1, \dots, y_n) = r(y_1, \dots, y_n) \cdot a_{y_n, y_{n+1}} = r(y_1, \dots, y_n) \cdot a_{y_n, \text{STOP}}$$

$$\pi(k, v) = \max_{(y_1, \dots, y_n) \in S(k, v)} r(y_1, \dots, y_k)$$

Where $S(k, v)$ is the set of tag sequences y_1, \dots, y_k such that $y_k = v$.

$$\pi(0, \text{START}) = 1$$

$$\pi(1, v) = \max_{u \in T} \{ \pi(0, \text{START}) \cdot a_{\text{START}, u} \cdot b_v(x_k) \}$$

The algorithm is initialized whereby for the first input token, we calculate the product of the transition probability from the 'START' tag to a possible tag and the emission probability of the same tag to the input token. This probability is calculated for every possible tag recursively:

$$\pi(k, v) = \max_{u \in T} \{ \pi(k-1, u) \cdot a_{u, v} \cdot b_v(x_k) \}$$

The next token in the sequence calculates the product of the transition and emission probabilities for each possible tag along with the total probability for the previous tag. This is repeated until the last word of sequence, where the calculation involves only the probability for the previous tag and the transition probability between the previous tag and the 'STOP' tag.

$$\max_{y_1, \dots, y_n} p(x_1, \dots, x_n, y_0 = \text{START}, y_1, \dots, y_n, y_{n+1} = \text{STOP}) = \max_{v \in T} \{ \pi(n, v) \cdot a_{v, \text{STOP}} \}$$

$$y_n^* = \arg \max_v \{ \pi(n, v) \cdot a_{v, \text{STOP}} \}$$

$$y_{n-1}^* = \arg \max_u \{ \pi(n-1, u) \cdot a_{u, y_n^*} \}$$

$$y_1^*, \dots, y_n^* = \arg \max_{y_1, \dots, y_n} p(x_1, \dots, x_n, y_1, \dots, y_n)$$

The Viterbi method takes the input file tokens, the transition parameters and the emission parameters and outputs a predicted tag sequence (for a given token sequence).

II. Result

	EN	FR	CN	SG
<i>Entity in gold data</i>	226	223	362	1382
<i>Entity in prediction</i>	162	166	158	717
<i>Correct entity</i>	104	112	63	382
<i>Entity precision</i>	0.6420	0.6747	0.3987	0.5328
<i>Entity recall</i>	0.4602	0.5022	0.1740	0.2764
<i>Entity F</i>	0.5361	0.5758	0.2423	0.3640
<i>Correct sentiment</i>	64	72	46	242
<i>Sentiment precision</i>	0.3951	0.4337	0.2911	0.3375
<i>Sentiment recall</i>	0.2832	0.3229	0.1271	0.1751
<i>Sentiment F</i>	0.3299	0.3702	0.1769	0.2306

Part 4: Max-Marginal Decoding Algorithm

I. Approach

The max-marginal decoding algorithm calculates the forward and backward scores for a given token at a given position in the sequence. The alpha or the forward scores, $\alpha_u(j)$ are calculated by taking the sum of scores of all paths from START to the state $\langle j, u \rangle$. The beta values or the backward scores, $\beta_u(j)$ are calculated by taking the sum of scores of all paths from state $\langle j, u \rangle$ to the final state STOP. For a given token sequence, the forward and backward probabilities are calculated recursively for all possible path sequences.

The optimal tag y at the position i given the token sequence, is found by computing the marginal probability term:

$$p(y_i = u \mid x_1, \dots, x_n; \theta) = \frac{p(x_1, x_2, \dots, x_{i-1}, y_i = u, x_i, \dots, x_n; \theta)}{p(x_1, \dots, x_n; \theta)} = \frac{p(x_1, x_2, \dots, x_{i-1}, y_i = u; \theta) p(x_i, \dots, x_n \mid y_i = u; \theta)}{p(x_1, \dots, x_n; \theta)}$$

$$= \frac{\alpha_u(i) \beta_u(i)}{\sum_v \alpha_v(j) \beta_v(j)}$$

$$y_i^* = \arg \max_u p(y_i = u \mid x_1, \dots, x_n; \theta) = \arg \max_u \frac{\alpha_u(i) \beta_u(i)}{\sum_v \alpha_v(j) \beta_v(j)} = \arg \max_u \alpha_u(i) \beta_u(i)$$

The maxMarginal method takes the input file tokens, the transition parameters and the emission parameters and outputs a predicted tag sequence (for a given token sequence). Unlike the previous algorithms, our maxMarginal method takes the input token sequences, sentence by sentence (instead of all the input file tokens).

II. Result

	EN	FR
<i>Entity in gold data</i>	226	223
<i>Entity in prediction</i>	303	296
<i>Correct entity</i>	134	171
<i>Entity precision</i>	0.4422	0.5777
<i>Entity recall</i>	0.5929	0.7668
<i>Entity F</i>	0.5066	0.6590
<i>Correct sentiment</i>	79	96
<i>Sentiment precision</i>	0.2607	0.3243
<i>Sentiment recall</i>	0.3496	0.4305
<i>Sentiment F</i>	0.2987	0.3699

Part 5: Challenge – More efficient handling of data

I. Approach

We wanted to try and improve the max marginal model by filtering and transforming the data so it would yield more accurate results. We do this by:

a. Non-Case Sensitive

As we believe that an entity and sentiment analysis is not case-sensitive, we decided to convert all the word to lowercase

b. Classification

We further improved the emission probabilities, by classifying the unknown values into unknown word classes:

1. Numeric: Unknown word containing at least one numeric character
2. Symbols: Unknown word containing at least one symbol (\$, (,), !, ., -, _, ', etc.)
3. Others: Unknown word that does not fit into the other classes

II. Result

	Previous implementation	
	EN	FR
<i>Entity in gold data</i>	226	223
<i>Entity in prediction</i>	303	296
<i>Correct entity</i>	134	171
<i>Entity precision</i>	0.4422	0.5777
<i>Entity recall</i>	0.5929	0.7668
<i>Entity F</i>	0.5066	0.6590
<i>Correct sentiment</i>	79	96
<i>Sentiment precision</i>	0.2607	0.3243
<i>Sentiment recall</i>	0.3496	0.4305
<i>Sentiment F</i>	0.2987	0.3699

	EN	FR
<i>Entity in gold data</i>	226	223
<i>Entity in prediction</i>	314	303
<i>Correct entity</i>	140	175
<i>Entity precision</i>	0.4459	0.5576
<i>Entity recall</i>	0.6195	0.7848
<i>Entity F</i>	0.5185	0.6654
<i>Correct sentiment</i>	81	110
<i>Sentiment precision</i>	0.2580	0.3630
<i>Sentiment recall</i>	0.3584	0.4933
<i>Sentiment F</i>	0.3000	0.4183

As can be seen, the model performs better on an average.

Calculating Emission Parameter

Each emission parameter $b_y(x) = p(X = x | Y = y)$ for $y = 1, \dots, |T| - 2$ and $x \in X$, is the probability of emitting word x given a tag y : $\sum_{x \in X} b_y(x) = 1$ for any y . The emission parameters are essentially calculated by counting the number of times a word x has a tag y , this result is then divided by the number of times tag y appears in the input file.

$$e(x|y) \rightarrow \frac{\text{Count}(y \rightarrow x)}{\text{Count}(y)}$$

The calculateEmission method takes the input file tags, tokens (words) and a learning threshold (to determine tokens to be classified as unknown) and outputs the emission parameters.

Calculating Transition Parameter

Each transition parameter $a_{i,j} = p(y_{next} = j | y_{current} = i)$ for $i = 0, 1, \dots, |T| - 2$ and $j = 1, \dots, |T| - 1$, is the probability of transitioning from state i to state j : $\sum_{k=1}^{|T|-1} a_{i,k} = 1$ for any i . The transition parameters are thus calculated by counting the number of times a tag y_{i-1} transitions to tag y_i , this result is divided by the number of times tag y_{i-1} appears in the input file.

$$q(y_i|y_{i-1}) \rightarrow \frac{\text{Count}(y_{i-1} \rightarrow y_i)}{\text{Count}(y_{i-1})}$$

The calculateTransition method takes the input file tags and outputs the transition parameters.

Miscellaneous

Parsing of Training Data

We separate the training data into two lists for training, one containing the input tokens and the other containing the input tags. This allows us to calculate the emission and transition parameters with greater ease.

Parsing of Input Data

The input data is stripped of white spaces and contains the input tokens (whose tag sequence is to be predicted).